

# ProjectDataAnalysis

March 9, 2025

## 1 Reading Data

Thyroid cancer data is retrieved from Kaggle website: <https://www.kaggle.com/datasets/ankushpanday1/thyroid-cancer-risk-prediction-dataset/data>

```
[ ]: #import necessary libraries
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import prince
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from scipy.stats import randint
from sklearn.decomposition import PCA
```

```
[42]: #Mount google drive
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[4]: #read data from drive
thyroid_data = pd.read_csv('/content/drive/My Drive/COURSES/Winter_2025/CIND820_
    XJH - Big Data Analytics Project - W202/thyroid_cancer_risk_data.csv')
print(thyroid_data.head())
```

	Patient_ID	Age	Gender	Country	Ethnicity	Family_History	\
0	1	66	Male	Russia	Caucasian	No	
1	2	29	Male	Germany	Hispanic	No	

2	3	86	Male	Nigeria	Caucasian	No
3	4	75	Female	India	Asian	No
4	5	35	Female	Germany	African	Yes

	Radiation_Exposure	Iodine_Deficiency	Smoking	Obesity	Diabetes	TSH_Level	\
0	Yes	No	No	No	No	9.37	
1	Yes	No	No	No	No	1.83	
2	No	No	No	No	No	6.26	
3	No	No	No	No	No	4.10	
4	Yes	No	No	No	No	9.10	

	T3_Level	T4_Level	Nodule_Size	Thyroid_Cancer_Risk	Diagnosis
0	1.67	6.16	1.08	Low	Benign
1	1.73	10.54	4.05	Low	Benign
2	2.59	10.57	4.61	Low	Benign
3	2.62	11.04	2.46	Medium	Benign
4	2.11	10.71	2.11	High	Benign

## 2 Preprocessing

Data is explored to know attributes name, shape of the data frame, data types, missing values and regional location of thyroid cancer patients.

## 3 Variable attributes, Data Types and Missing Values

```
[6]: # Display the column names of the DataFrame
print("Column names in the DataFrame:")
print(thyroid_data.columns)

#Display the shape of the DataFrame
print("Shape of the DataFrame:")
print(thyroid_data.shape)

#Display the data types of the columns
print("Data types of the columns:")
print(thyroid_data.dtypes)

#Display the missing values in the DataFrame
missing_values = thyroid_data.isnull().sum()
print("Missing values in the DataFrame:", missing_values)

#Different countries in the dataset
countries = thyroid_data['Country'].unique()
print("Countries in the dataset:", countries)
```

Column names in the DataFrame:  
Index(['Patient\_ID', 'Age', 'Gender', 'Country', 'Ethnicity', 'Family\_History',

```

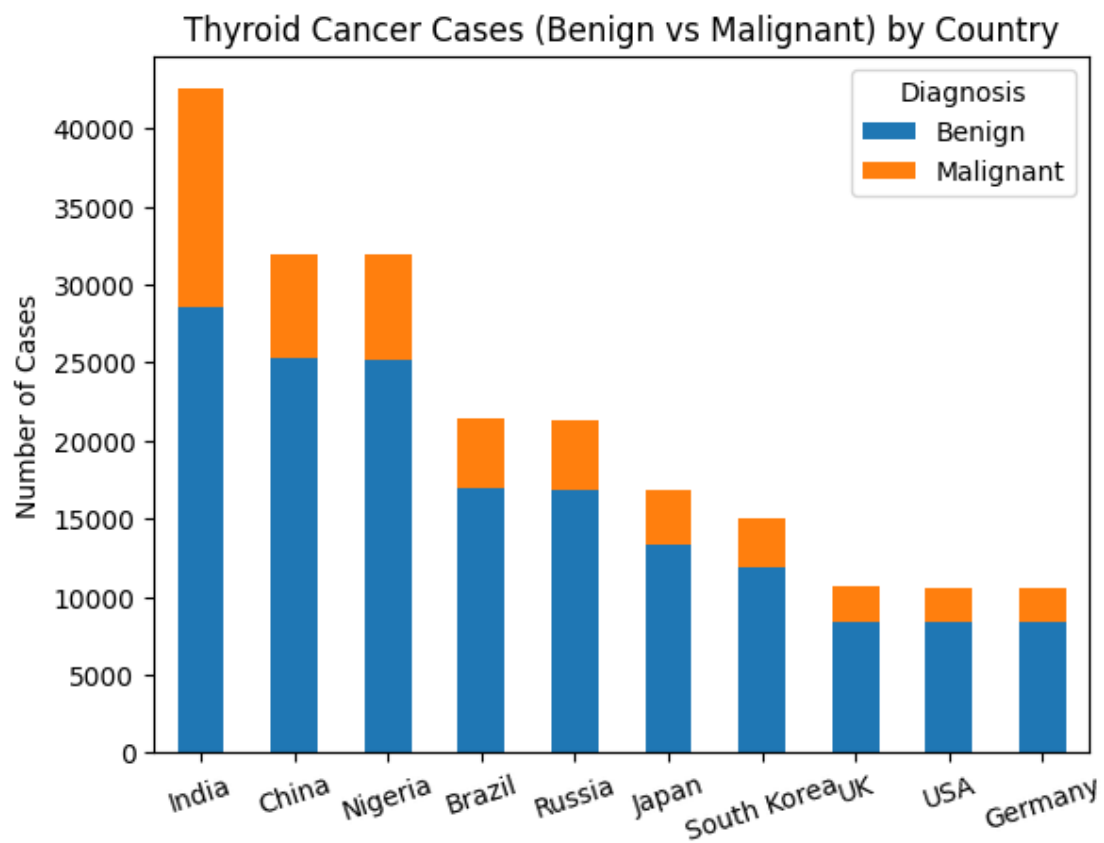
        'Radiation_Exposure', 'Iodine_Deficiency', 'Smoking', 'Obesity',
        'Diabetes', 'TSH_Level', 'T3_Level', 'T4_Level', 'Nodule_Size',
        'Thyroid_Cancer_Risk', 'Diagnosis'],
        dtype='object')
Shape of the DataFrame:
(212691, 17)
Data types of the columns:
Patient_ID          int64
Age                 int64
Gender              object
Country             object
Ethnicity           object
Family_History      object
Radiation_Exposure  object
Iodine_Deficiency   object
Smoking             object
Obesity             object
Diabetes            object
TSH_Level           float64
T3_Level            float64
T4_Level            float64
Nodule_Size         float64
Thyroid_Cancer_Risk object
Diagnosis           object
dtype: object
Missing values in the DataFrame: Patient_ID          0
Age          0
Gender       0
Country      0
Ethnicity    0
Family_History 0
Radiation_Exposure 0
Iodine_Deficiency 0
Smoking       0
Obesity       0
Diabetes      0
TSH_Level     0
T3_Level      0
T4_Level      0
Nodule_Size   0
Thyroid_Cancer_Risk 0
Diagnosis     0
dtype: int64
Countries in the dataset: ['Russia' 'Germany' 'Nigeria' 'India' 'UK' 'South
Korea' 'Brazil' 'China'
'Japan' 'USA']

```

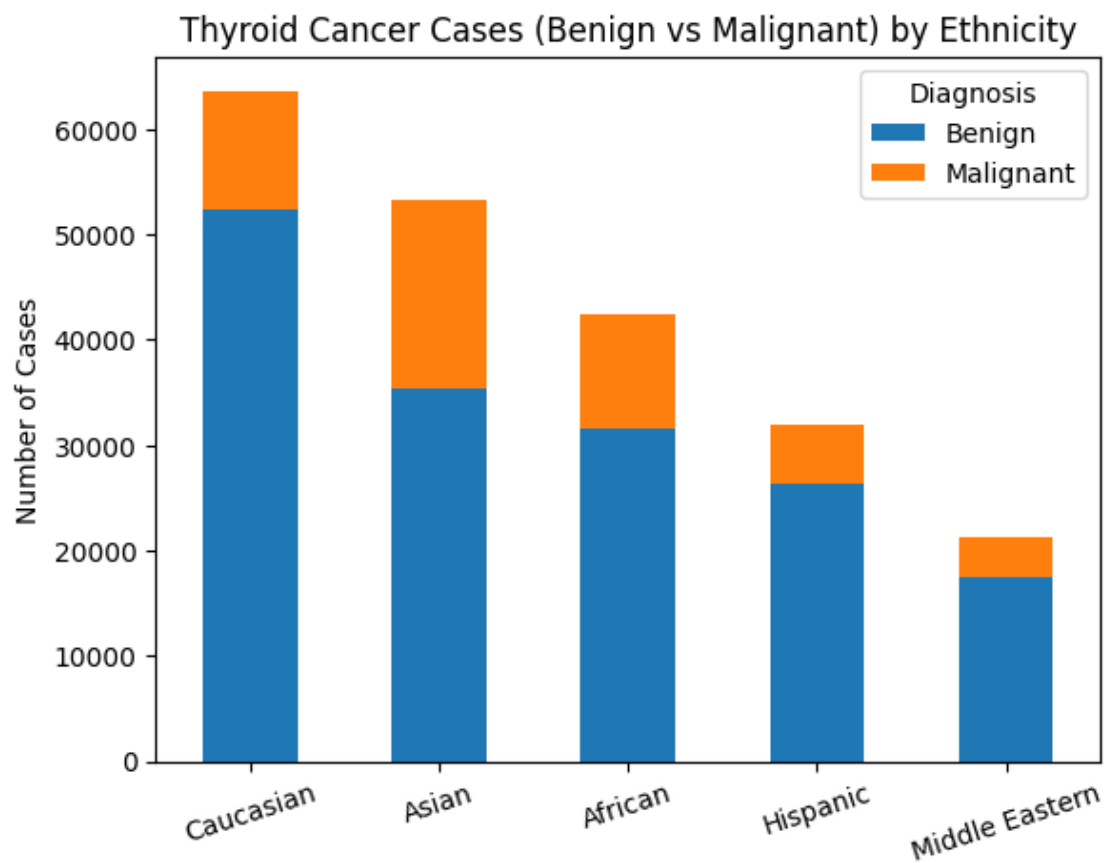
## 4 Thyroid Cancer Distribution by Region, Gender and Ethnicity

```
[41]: def myPlot(data, var):  
    plt.figure(figsize=(6, 5))  
    sorted_data = thyroid_data.groupby([var, 'Diagnosis']).size().unstack().  
    ↪fillna(0)  
    sorted_data['Total'] = sorted_data.sum(axis=1)  
    sorted_data = sorted_data.sort_values(by='Total', ascending=False).  
    ↪drop(columns='Total')  
    sorted_data.plot(kind='bar', stacked=True)  
    plt.title('Thyroid Cancer Cases (Benign vs Malignant) by '+var)  
    plt.xlabel(' ')  
    plt.ylabel('Number of Cases')  
    plt.xticks(rotation=18)  
    plt.legend(title='Diagnosis')  
    plt.show()  
  
myPlot(thyroid_data, 'Country')  
myPlot(thyroid_data, 'Ethnicity')  
myPlot(thyroid_data, 'Gender')
```

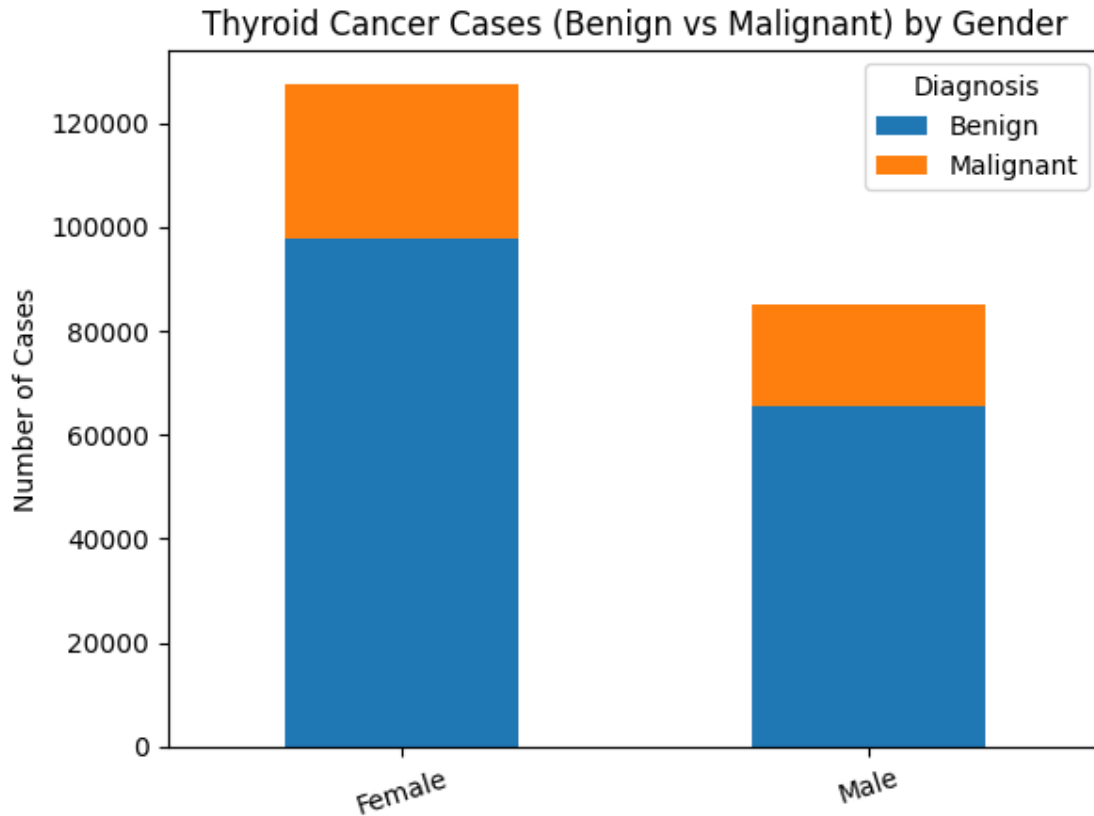
<Figure size 600x500 with 0 Axes>



<Figure size 600x500 with 0 Axes>



<Figure size 600x500 with 0 Axes>



## 5 Summary Statistics

Proportion by categorical variables rounded to 2 decimal places

```
[10]: print("Proportion by categorical variables (rounded to 2 decimal places):")
categorical_columns = thyroid_data.select_dtypes(include=['object']).columns
for column in categorical_columns:
    print(f"\n{column}: \n{thyroid_data[column].value_counts(normalize=True)}.
    ↪round(2)*100}")
```

Proportion by categorical variables (rounded to 2 decimal places):

Gender:

Gender

Female 60.0

Male 40.0

Name: proportion, dtype: float64

Country:

Country

India	20.0
China	15.0
Nigeria	15.0
Brazil	10.0
Russia	10.0
Japan	8.0
South Korea	7.0
UK	5.0
USA	5.0
Germany	5.0

Name: proportion, dtype: float64

Ethnicity:

Ethnicity	
Caucasian	30.0
Asian	25.0
African	20.0
Hispanic	15.0
Middle Eastern	10.0

Name: proportion, dtype: float64

Family\_History:

Family_History	
No	70.0
Yes	30.0

Name: proportion, dtype: float64

Radiation\_Exposure:

Radiation_Exposure	
No	85.0
Yes	15.0

Name: proportion, dtype: float64

Iodine\_Deficiency:

Iodine_Deficiency	
No	75.0
Yes	25.0

Name: proportion, dtype: float64

Smoking:

Smoking	
No	80.0
Yes	20.0

Name: proportion, dtype: float64

Obesity:

Obesity	
---------	--



```
No      70.0
Yes      30.0
Name: proportion, dtype: float64
```

```
Diabetes:
Diabetes
No      80.0
Yes      20.0
Name: proportion, dtype: float64
```

```
Thyroid_Cancer_Risk:
Thyroid_Cancer_Risk
Low      51.0
Medium   34.0
High     15.0
Name: proportion, dtype: float64
```

```
Diagnosis:
Diagnosis
Benign      77.0
Malignant   23.0
Name: proportion, dtype: float64
```

Summary statistics of all continuous variables rounded to 2 decimal places

```
[6]: continuous_columns = thyroid_data.select_dtypes(include=['number']).columns
continuous_columns_except_id = continuous_columns.drop('Patient_ID',
↳errors='ignore')
summary_stats = thyroid_data[continuous_columns_except_id].describe().round(2)
print(summary_stats)
```

	Age	TSH_Level	T3_Level	T4_Level	Nodule_Size
count	212691.00	212691.00	212691.00	212691.00	212691.00
mean	51.92	5.05	2.00	8.25	2.50
std	21.63	2.86	0.87	2.16	1.44
min	15.00	0.10	0.50	4.50	0.00
25%	33.00	2.57	1.25	6.37	1.25
50%	52.00	5.04	2.00	8.24	2.51
75%	71.00	7.52	2.75	10.12	3.76
max	89.00	10.00	3.50	12.00	5.00

Scaling continuous variables

```
[7]: # Initialize the scaler
scaler = MinMaxScaler()

# Fit and transform the continuous variables
scaled_df = scaler.fit_transform(thyroid_data[continuous_columns_except_id])
```

```
# Create a new DataFrame with the scaled data
scaled_data = pd.DataFrame(scaled_df, columns=continuous_columns_except_id)

# Print the scaled data
print(scaled_data.head())
```

	Age	TSH_Level	T3_Level	T4_Level	Nodule_Size
0	0.689189	0.936364	0.390000	0.221333	0.216
1	0.189189	0.174747	0.410000	0.805333	0.810
2	0.959459	0.622222	0.696667	0.809333	0.922
3	0.810811	0.404040	0.706667	0.872000	0.492
4	0.270270	0.909091	0.536667	0.828000	0.422

Merge the scaled data back into the original Thyroid data frame.

```
[8]: thyroid_data = pd.concat([thyroid_data,
    ↪ drop(columns=continuous_columns_except_id), scaled_data], axis=1)

# Print the head of the merged DataFrame
print(thyroid_data.head())
```

	Patient_ID	Gender	Country	Ethnicity	Family_History	Radiation_Exposure	\
0	1	Male	Russia	Caucasian	No	Yes	
1	2	Male	Germany	Hispanic	No	Yes	
2	3	Male	Nigeria	Caucasian	No	No	
3	4	Female	India	Asian	No	No	
4	5	Female	Germany	African	Yes	Yes	

	Iodine_Deficiency	Smoking	Obesity	Diabetes	Thyroid_Cancer_Risk	Diagnosis	\
0	No	No	No	No	Low	Benign	
1	No	No	No	No	Low	Benign	
2	No	No	No	No	Low	Benign	
3	No	No	No	No	Medium	Benign	
4	No	No	No	No	High	Benign	

	Age	TSH_Level	T3_Level	T4_Level	Nodule_Size
0	0.689189	0.936364	0.390000	0.221333	0.216
1	0.189189	0.174747	0.410000	0.805333	0.810
2	0.959459	0.622222	0.696667	0.809333	0.922
3	0.810811	0.404040	0.706667	0.872000	0.492
4	0.270270	0.909091	0.536667	0.828000	0.422

## 6 Predictive Analysis

The structure of predictive analysis - first we split the data as train and test, then choose a model, training the model, evaluate the model, and deploy the model for prediction.

## 7 Split the data into training and testing sets

```
[9]: X = thyroid_data.drop('Diagnosis', axis=1) # Features
y = thyroid_data['Diagnosis'] # Target variable

# Convert categorical features to numerical using one-hot encoding
X = pd.get_dummies(X, drop_first=True)

# random_state for reproducibility
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (170152, 28)
X_test shape: (42539, 28)
y_train shape: (170152,)
y_test shape: (42539,)
```

## 8 K-Nearest Neighbour

Grid search automates the process of testing multiple K values along with other hyperparameters. We can use GridSearchCV from scikit-learn libraries to search through different combinations of K values and evaluate the performance using cross-validation. The optimal value of k is found to be 5.

```
[34]: # Use the optimal k found from GridSearchCV
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
```

```
Accuracy: 0.7246526716660006
```

	precision	recall	f1-score	support
Benign	0.77	0.92	0.84	32615
Malignant	0.25	0.09	0.13	9924
accuracy			0.72	42539
macro avg	0.51	0.50	0.48	42539

weighted avg	0.65	0.72	0.67	42539
--------------	------	------	------	-------

```
[[29935 2680]
 [ 9033  891]]
```

## 9 Logistic Regression

```
[13]: # Initialize and train the logistic regression model
logreg = LogisticRegression(max_iter=1000)
logreg.fit(X_train, y_train)

# Make predictions on the test set
y_pred_logreg = logreg.predict(X_test)

# Evaluate the model
accuracy_logreg = accuracy_score(y_test, y_pred_logreg)
print(f"Logistic Regression Accuracy: {accuracy_logreg}")
print(classification_report(y_test, y_pred_logreg))
print(confusion_matrix(y_test, y_pred_logreg))
```

```
Logistic Regression Accuracy: 0.8250546557276852
      precision    recall  f1-score   support

   Benign       0.85      0.94      0.89      32615
  Malignant       0.69      0.45      0.54       9924

 accuracy                   0.83      42539
 macro avg       0.77      0.69      0.72      42539
weighted avg       0.81      0.83      0.81      42539

[[30657 1958]
 [ 5484 4440]]
```

## 10 Decision Tree

The optimization of decision tree classifier is performed by pre-pruning; maximum depth of the tree is used as a control variable for pre-pruning. Also, the attribute selection measure ‘entropy’ for the information gain is used as control parameter.

```
[16]: # Create Decision Tree classifier object
dtree = DecisionTreeClassifier(criterion="entropy", max_depth=5)

# Train Decision Tree Classifier
dtree.fit(X_train,y_train)

#Predict the response for test dataset
y_pred_dtree = dtree.predict(X_test)
```

```
# Evaluate the model
accuracy_dtree = accuracy_score(y_test, y_pred_dtree)
print(f"Decision Tree Accuracy: {accuracy_dtree}")
print(classification_report(y_test, y_pred_dtree))
print(confusion_matrix(y_test, y_pred_dtree))
```

```
Decision Tree Accuracy: 0.825007640047956
      precision    recall  f1-score   support

   Benign       0.85      0.94      0.89      32615
  Malignant       0.69      0.45      0.54       9924

 accuracy
macro avg       0.77      0.69      0.72      42539
weighted avg     0.81      0.83      0.81      42539

[[30654  1961]
 [ 5483  4441]]
```

## 11 Random Forest

Create an instance of the Random forest model with the default parameters. Then fit the model to training data; pass both the features and the target variable so the model can learn to predict.

```
[17]: # Train the Random Forest Classifier
rf_classifier = RandomForestClassifier()
rf_classifier.fit(X_train, y_train)

# Make predictions on the test set
y_pred_rf = rf_classifier.predict(X_test)

# Evaluate the model
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print(f"Random Forest Accuracy: {accuracy_rf}")
print(classification_report(y_test, y_pred_rf))
print(confusion_matrix(y_test, y_pred_rf))
```

```
Random Forest Accuracy: 0.8243024048520181
      precision    recall  f1-score   support

   Benign       0.85      0.94      0.89      32615
  Malignant       0.69      0.44      0.54       9924

 accuracy
macro avg       0.77      0.69      0.72      42539
weighted avg     0.81      0.82      0.81      42539
```

```
[[30685 1930]
 [ 5544 4380]]
```

## 12 Principal Component Analysis

Since the dataset has both the continuous and categorical feature variables; the principal component analysis will be based on the FAMD() function from Python library Prince. The FAMD () function implements the principal component method dedicated to explore data with both continuous and categorical variables.

```
[27]: # prompt: PCA using prince library
pca = prince.PCA(n_components=2) # Reduce to 2 principal components
pca = pca.fit(X_train)
X_train_pca = pca.transform(X_train)
X_test_pca = pca.transform(X_test)

# You can now use X_train_pca in your models
print(X_train_pca.head())
```

```
component      0      1
117332    -1.145959 -0.074020
63420      0.426001  1.437753
179948      0.572743  2.263643
187371     -1.081722  1.805283
103129      0.221810  0.270046
```

Prediction accuracy with KNN after PCA

```
[29]: # Example with KNN after PCA
knn_pca = KNeighborsClassifier(n_neighbors=5)
knn_pca.fit(X_train_pca, y_train)
y_pred_knn_pca = knn_pca.predict(X_test_pca)
accuracy_knn_pca = accuracy_score(y_test, y_pred_knn_pca)
print(f"KNN Accuracy with PCA: {accuracy_knn_pca}")
print(classification_report(y_test, y_pred_knn_pca))
print(confusion_matrix(y_test, y_pred_knn_pca))
```

KNN Accuracy with PCA: 0.7991020005171725

	precision	recall	f1-score	support
Benign	0.83	0.92	0.88	32615
Malignant	0.61	0.39	0.48	9924
accuracy			0.80	42539
macro avg	0.72	0.66	0.68	42539
weighted avg	0.78	0.80	0.78	42539

```
[[30124 2491]
 [ 6055 3869]]
```

Prediction accuracy with decision Tree after PCA

```
[30]: # Example with Decision Tree after PCA
dtree_pca = DecisionTreeClassifier(criterion="entropy", max_depth=5)
dtree_pca.fit(X_train_pca, y_train)
y_pred_dtree_pca = dtree_pca.predict(X_test_pca)
accuracy_dtree_pca = accuracy_score(y_test, y_pred_dtree_pca)
print(f"Decision Tree Accuracy with PCA: {accuracy_dtree_pca}")
print(classification_report(y_test, y_pred_dtree_pca))
print(confusion_matrix(y_test, y_pred_dtree_pca))
```

Decision Tree Accuracy with PCA: 0.8249371165283622

	precision	recall	f1-score	support
Benign	0.85	0.94	0.89	32615
Malignant	0.69	0.45	0.54	9924
accuracy			0.82	42539
macro avg	0.77	0.69	0.72	42539
weighted avg	0.81	0.82	0.81	42539

```
[[30654 1961]
 [ 5486 4438]]
```

Prediction accuracy with Random Forest after PCA

```
[31]: # Example with Random Forest after PCA
rf_pca = RandomForestClassifier()
rf_pca.fit(X_train_pca, y_train)
y_pred_rf_pca = rf_pca.predict(X_test_pca)
accuracy_rf_pca = accuracy_score(y_test, y_pred_rf_pca)
print(f"Random Forest Accuracy with PCA: {accuracy_rf_pca}")
print(classification_report(y_test, y_pred_rf_pca))
print(confusion_matrix(y_test, y_pred_rf_pca))
```

Random Forest Accuracy with PCA: 0.8042737252873834

	precision	recall	f1-score	support
Benign	0.83	0.93	0.88	32615
Malignant	0.63	0.39	0.48	9924
accuracy			0.80	42539
macro avg	0.73	0.66	0.68	42539
weighted avg	0.79	0.80	0.79	42539

```
[[30319 2296]
```

```
[ 6030  3894]]
```

## 13 Model Performance Comparison

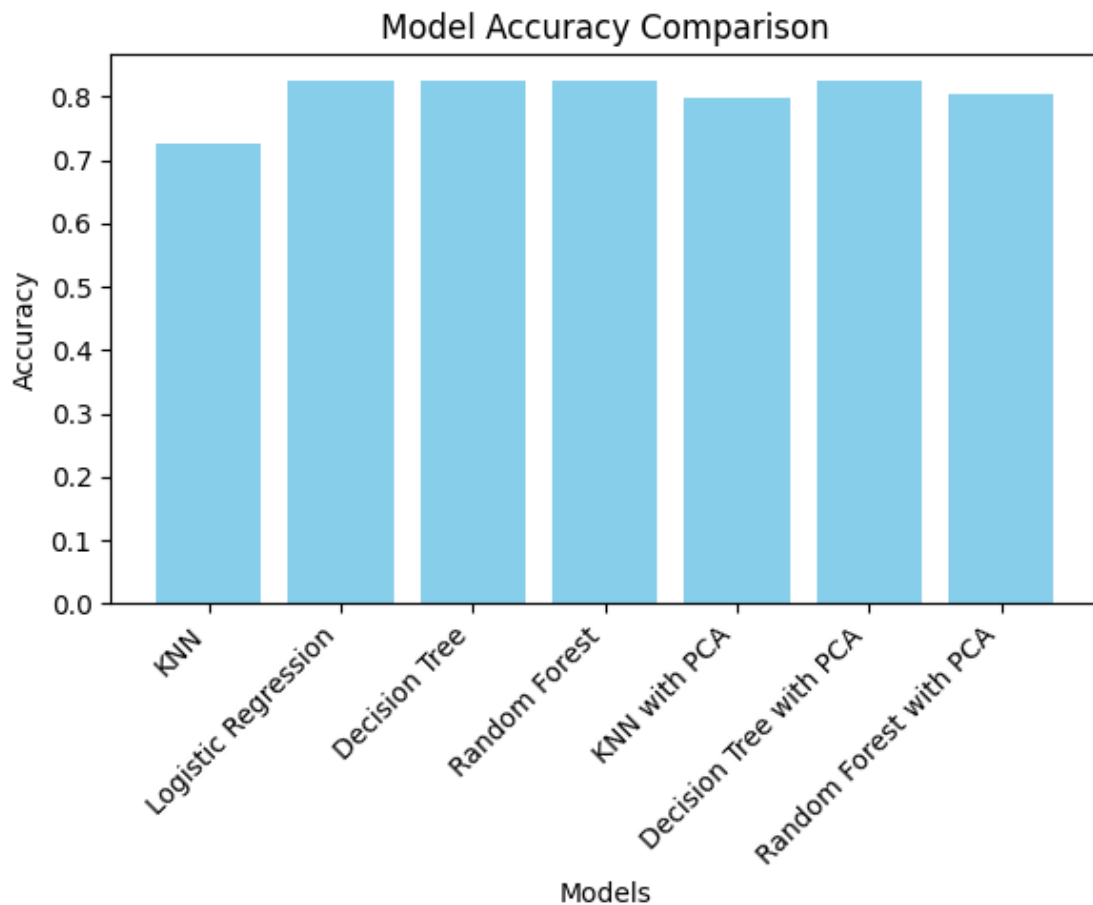
Make a bar chart of all above accuracies found from different machine learning classifiers.

```
[39]: accuracies = {
    'KNN': accuracy,
    'Logistic Regression': accuracy_logreg,
    'Decision Tree': accuracy_dtree,
    'Random Forest': accuracy_rf,
    'KNN with PCA': accuracy_knn_pca,
    'Decision Tree with PCA': accuracy_dtree_pca,
    'Random Forest with PCA': accuracy_rf_pca
}

models = list(accuracies.keys())
accuracy_values = list(accuracies.values())

plt.figure(figsize=(6, 5))
plt.bar(models, accuracy_values, color='skyblue')
plt.xlabel("Models")
plt.ylabel("Accuracy")
plt.title("Model Accuracy Comparison")
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for better
    ↪ readability
plt.tight_layout() # Adjust layout to prevent labels from overlapping
plt.show()
```





Model accuracies in data frame view.

```
[40]: df = pd.DataFrame(list(accuracies.items()), columns=['Model', 'Accuracy'])  
df
```

```
[40]:
```

	Model	Accuracy
0	KNN	0.724653
1	Logistic Regression	0.825055
2	Decision Tree	0.825008
3	Random Forest	0.824302
4	KNN with PCA	0.799102
5	Decision Tree with PCA	0.824937
6	Random Forest with PCA	0.804274